

SIM-OPS: An Autonomous Multi-Agent Framework for Intelligent Business Operations Management

1st Arshvir Singh Kalsi

*Computer Engineering Department
St. Francis Institute of Technology
Mumbai, India
arshvirsk26@student.sfit.ac.in*

2nd Viraj Prabhu

*Computer Engineering Department
St. Francis Institute of Technology
Mumbai, India
virajprabhu47@student.sfit.ac.in*

3rd Ghrani Poojari

*Computer Engineering Department
St. Francis Institute of Technology
Mumbai, India
ghranipoojari1912@student.sfit.ac.in*

4th Hasti Shah

*Computer Engineering Department
St. Francis Institute of Technology
Mumbai, India
shahhasti27@student.sfit.ac.in*

5th Varsha Nagpurkar

*Computer Engineering Department
St. Francis Institute of Technology
Mumbai, India
varshanagpurkar@sfit.ac.in*

Abstract—This paper presents SIM-OPS (Smart Intelligence Management for Operations Planning & Supervision), an autonomous multi-agent system designed to function as an intelligent business operations manager. The system leverages advanced machine learning models (XGBoost, Isolation Forests), real-time data processing, and coordinated agent orchestration powered by Large Language Models (LLMs) to continuously monitor organizational metrics, predict business risks, and execute automated interventions. Our framework integrates six specialized agents—Monitoring, Prediction, Decision, Action, Reporting, and Feedback—that communicate through a message bus to perform end-to-end business intelligence operations. The Decision engine utilizes the ReAct (Reasoning and Acting) paradigm to synthesize context-aware retention strategies. The system achieves 86.94% accuracy in churn prediction, processes predictions autonomously every 6 hours, and executes multi-channel alerts through Slack, Jira, and email. Experimental results demonstrate that the autonomous agent chain reduces manual intervention by 90% while maintaining 96% decision accuracy and 0.3-second average response time for critical alerts. The system is deployed as a production-ready web application built with Next.js 14, TypeScript, Supabase, and LangChain, with a Python-based FastAPI machine learning service.

Index Terms—Multi-agent systems, autonomous operations, machine learning, business intelligence, churn prediction, workflow automation, real-time monitoring

I. INTRODUCTION

A. Motivation

Modern businesses generate vast amounts of operational data from customer interactions, financial transactions, support systems, and product usage. However, extracting actionable insights from this data and executing timely interventions remains a significant challenge. Traditional business intelligence systems require constant human oversight, leading to delayed

responses to critical events such as customer churn, revenue anomalies, and operational inefficiencies.

The need for autonomous, intelligent systems that can continuously observe, analyze, and act on business data has become increasingly critical. Such systems must not only detect patterns and predict outcomes but also coordinate multiple specialized components to execute complex workflows without human intervention.

B. Contributions

This paper makes the following contributions:

- **Multi-Agent Architecture:** We present a novel six-agent framework where specialized agents collaborate through a message bus to perform comprehensive business operations management.
- **Autonomous Prediction Pipeline:** A fully automated system that runs ML predictions every 6 hours, processes customer data, and triggers agent chains based on risk thresholds.
- **Intelligent Decision Engine:** A rule-based decision system that evaluates predictions, applies business logic, and determines appropriate actions with 94% confidence.
- **Multi-Channel Action Execution:** An action executor that delivers alerts through Slack, creates Jira tickets, and sends emails based on severity classification.
- **Production-Ready Implementation:** A complete system deployed with modern web technologies, achieving sub-second response times and 24/7 autonomous operation.

C. Paper Organization

The remainder of this paper is organized as follows: Section II reviews related work in multi-agent systems and business intelligence. Section III describes the system architecture and

agent design. Section IV details the machine learning models and data pipeline. Section V presents the implementation and deployment. Section VI provides experimental results and performance analysis. Section VII discusses limitations and future work. Section VIII concludes the paper.

II. RELATED WORK

A. Multi-Agent Systems

Multi-agent systems (MAS) have been extensively studied in artificial intelligence research [1]. Traditional MAS frameworks focus on agent communication protocols, coordination mechanisms, and distributed problem-solving. However, most existing systems operate in simulated environments or require significant human oversight.

Recent work in autonomous agents has explored reinforcement learning approaches [2] and hierarchical agent architectures [3]. Our work differs by focusing on practical business operations with real-world constraints, emphasizing reliability, explainability, and integration with existing enterprise systems.

B. Business Intelligence and Predictive Analytics

Predictive analytics in business intelligence has traditionally relied on batch processing and periodic reporting [4]. Modern approaches incorporate real-time streaming analytics [5] and automated machine learning (AutoML) [6].

Customer churn prediction has been addressed using various machine learning techniques including logistic regression [7], random forests [8], and gradient boosting [9]. Our system extends these approaches by integrating predictions into an autonomous action execution framework.

C. Workflow Automation

Workflow automation systems like Apache Airflow [10] and Temporal [11] provide orchestration capabilities but require manual workflow definition and lack intelligent decision-making. Business process management (BPM) systems [12] offer structured workflow execution but do not incorporate machine learning or autonomous agent coordination.

Our system combines workflow automation with intelligent agents that can dynamically determine appropriate actions based on real-time predictions and business rules.

III. THEORETICAL BACKGROUND

To provide a robust foundation for SIM-OPS, we synthesize principles from predictive machine learning, anomaly detection, and modern natural language processing.

A. Predictive Modeling with XGBoost

The core churn prediction model utilizes XGBoost (eXtreme Gradient Boosting), an optimized distributed gradient boosting library. XGBoost minimizes a regularized objective function:

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t) \quad (1)$$

where l is a differentiable convex loss function that measures the difference between the prediction $\hat{y}_i$ and the target y_i . The term Ω penalizes the complexity of the model:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2 \quad (2)$$

with T being the number of leaves in the tree and w the leaf weights. This regularization helps prevent overfitting, a common issue in customer behavioral datasets with high dimensionality.

B. Anomaly Detection via Isolation Forests

For monitoring KPIs such as revenue drops or sudden spikes in support tickets, we employ Isolation Forests. Unlike traditional methods that profile normal points, Isolation Forests explicitly isolate anomalies. The path length $h(x)$ to isolate a point x is equivalent to the number of edges from the root node to a terminating node. The anomaly score $s(x, n)$ is defined as:

$$s(x, n) = 2^{-\frac{E(h(x))}{c(n)}} \quad (3)$$

where $E(h(x))$ is the average path length and $c(n)$ is the average path length of unsuccessful searches in a Binary Search Tree. Anomalies are identified when $s(x, n)$ approaches 1, indicating that the point is easily isolated and thus abnormal.

C. LLM-based Agentic Reasoning

Our multi-agent orchestration relies on Large Language Models (LLMs) employing the ReAct (Reasoning and Acting) framework [13]. ReAct prompts the LLM to generate both reasoning traces and task-specific actions in an interleaved manner. Given a context c , the agent generates a sequence of thoughts t_i , actions a_i , and observations o_i :

$$\tau_i = \text{LLM}(c, t_1, a_1, o_1, \dots, t_{i-1}, a_{i-1}, o_{i-1}) \quad (4)$$

This paradigm allows the Decision Agent to dynamically evaluate the context of a high-risk customer and synthesize a tailored retention strategy rather than relying solely on static rule-based systems.

IV. SYSTEM ARCHITECTURE

A. Overview

SIM-OPS employs a layered architecture consisting of four primary components: (1) Data Ingestion Layer, (2) Machine Learning Service, (3) Multi-Agent Orchestration Layer, and (4) Action Execution Layer. Figure 1 illustrates the system architecture.

B. Data Ingestion Layer

The data ingestion layer collects data from multiple sources:

- **Stripe Webhooks:** Real-time payment events including successful transactions and payment failures
- **Customer Database:** User profiles, subscription data, and behavioral metrics stored in PostgreSQL (Supabase)
- **Analytics Events:** Product usage, feature adoption, and engagement metrics




Fig. 1. SIM-OPS System Architecture

- **Support Systems:** Ticket counts, resolution times, and customer satisfaction scores

Webhook handlers process incoming events and update the customer database in real-time, ensuring that ML models always operate on current data.

C. Machine Learning Service

The ML service is implemented as a Python FastAPI application providing four primary models:

1) *Churn Prediction Model:* Uses XGBoost classifier trained on customer features:

- Tenure (days since signup)
- Support ticket count
- Payment failure history
- Feature usage rate
- Session frequency and duration
- Total spend and discount usage

The model achieves 86.94% accuracy with 0.92 confidence on test data.

2) *Customer Lifetime Value (CLV) Model:* Predicts expected revenue per customer using gradient boosting regression on:

- Purchase history
- Average order value
- Purchase frequency
- Engagement score
- Referral activity

3) *Anomaly Detection Model:* Implements Isolation Forest algorithm to detect unusual patterns in time-series metrics such as revenue, user growth, and system costs.

4) *Revenue Forecasting Model:* Uses Prophet time-series model to generate 30-day revenue forecasts with confidence intervals.

All models are pre-trained and stored as pickle files, loaded on-demand for inference. The service exposes RESTful API endpoints for predictions.

D. Multi-Agent Orchestration Layer

The core innovation of SIM-OPS is its multi-agent orchestration system consisting of six specialized agents. The




Fig. 2. Predictive Analytics Interface displaying Machine Learning predictions and Churn risk distribution.

live operational status, active workflows, and agent execution statistics are monitored via a unified dashboard as shown in Figure 3.




Fig. 3. Autonomous Multi-Agent Monitoring and Configuration Portal.

1) *Monitoring Agent:* **Role:** KPI deviation detection and threshold monitoring

Responsibilities:

- Continuously monitors business metrics (churn rate, revenue, user growth)
- Compares current values against configured thresholds
- Detects anomalies using statistical methods
- Triggers alerts when thresholds are exceeded

Configuration:

```

{
  "churnRate": 3.0,
  "revenueDeviation": 5.0,
  "userGrowth": -2.0
}
  
```

2) *Prediction Agent:* **Role:** ML inference and risk scoring

Responsibilities:

- Fetches customer data from database
- Extracts features for ML models
- Executes churn and CLV predictions
- Stores predictions with confidence scores
- Identifies high-risk customers (>80% churn probability)

Performance: Processes 100 customers in 2.1 seconds average.

3) *Decision Agent:* **Role:** Severity classification and rule engine

Responsibilities:

- Receives inputs from Monitoring and Prediction agents
- Applies business rules to classify severity
- Determines appropriate actions based on risk level
- Routes decisions to Action Agent

Decision Rules:

- **Critical:** Churn $\geq 70\%$ AND worsening trend
- **High:** Churn $\geq 70\%$ OR multiple anomalies
- **Medium:** Churn $\geq 60\%$ AND worsening trend
- **Low:** All other cases

Figure 4 illustrates the Customer Detail view integrated with the Decision Agent's reasoning panel, displaying contributing features alongside the generated retention strategy.




Fig. 4. Customer Risk Profile and LLM-based Decision Reasoning Interface.

4) **Action Agent: Role:** Workflow execution and automation triggers

Responsibilities:

- Executes multi-channel notifications
- Creates Jira tickets for retention/operations teams
- Sends emails to account managers
- Triggers automated workflows
- Logs all execution results

Action Mapping:

- Critical: Slack + Jira + Email
- High: Slack + Jira
- Medium: Slack only
- Low: Database log only

5) **Reporting Agent: Role:** Summary generation and audit logging

Responsibilities:

- Aggregates activity from all agents
- Generates daily and weekly summaries
- Maintains comprehensive audit trails
- Delivers executive reports via email

6) **Feedback Agent: Role:** Outcome tracking and retraining triggers

Responsibilities:

- Analyzes model performance over time
- Detects model drift (confidence degradation)
- Tracks action success rates
- Triggers model retraining when needed

Evaluation Metrics:

- Average prediction confidence

- Confidence drift (prediction vs decision)
- Action success rate
- 7-day rolling window analysis

E. Agent Communication Protocol

Agents communicate through a message bus implemented using database-backed message passing. Each message contains:

```
{
  "from": "agent_id",
  "to": "agent_id",
  "type": "message_type",
  "payload": { ... },
  "timestamp": "ISO8601"
}
```

Messages are stored in the `agent_communications` table, providing full traceability and enabling asynchronous processing.

F. Action Execution Layer

The action execution layer implements integrations with external systems:

1) **Slack Integration:** Uses webhook API to send formatted alerts with color-coded severity indicators. Alerts include customer details, risk scores, and recommended actions.

2) **Jira Integration:** Creates tickets via REST API with appropriate project, issue type, and priority. Ticket descriptions include contributing factors and recommended actions.

3) **Email Integration:** Uses Resend API to send HTML-formatted emails to account managers and executives. Templates are dynamically rendered based on alert type.

G. LLM Orchestration with LangChain

To manage the interactions between the intelligent agents and the underlying LLM API, we utilize the LangChain framework. LangChain provides a standardized interface for memory management, prompt templating, and tool binding. The Decision Agent is implemented as a ReAct agent within LangChain, equipped with tools to query the CRM, fetch recent support interactions, and trigger the Action Agent.

The prompt template incorporates dynamic variables to provide context to the LLM:

```
PROMPT = """
Analyze the following customer risk profile:
Churn Probability: {churn_prob}
CLV: ${clv}
Recent Anomalies: {anomalies}
Recent Support Tickets: {support_history}
```

```
Formulate a targeted retention strategy using the
available tools. Ensure your reasoning is
documented before acting.
"""
```

This orchestration layer abstracts the complexity of API interactions and ensures that agent responses are parsed robustly into actionable JSON schemas before being dispatched to the Action Execution Layer.

V. IMPLEMENTATION

A. Technology Stack

1) Frontend and Backend:

- **Framework:** Next.js 14 with App Router
- **Language:** TypeScript for type safety
- **Styling:** Tailwind CSS v4 with shadcn/ui components
- **State Management:** TanStack React Query for server state
- **Charts:** Recharts for data visualization

The resulting user interface is a responsive, real-time dashboard displaying key business metrics, recent activities, and agent status logs (see Figure 5).




Fig. 5. SIM-OPS Real-Time Operations and KPI Monitoring Dashboard.

2) Database:

- **Database:** PostgreSQL via Supabase
- **ORM:** Supabase client with TypeScript types
- **Real-time:** Supabase Realtime for live updates
- **Authentication:** Supabase Auth with email/password

3) Machine Learning Service:

- **Framework:** FastAPI for high-performance API
- **ML Libraries:** scikit-learn, XGBoost, LightGBM
- **Time Series:** Prophet for forecasting
- **Data Processing:** NumPy, Pandas

B. Database Schema

The system uses 10 primary tables:

TABLE I
DATABASE SCHEMA OVERVIEW

Table	Rows	Purpose
agents	6	Agent metadata and status
agent_configs	6	Agent configuration and thresholds
agent_decisions	1000+	Decision history with reasoning
agent_metrics	6	Performance metrics per agent
agent_communications	500+	Inter-agent messages
workflows	3	Workflow definitions
workflow_runs	100+	Workflow execution logs
ml_predictions	5000+	ML prediction results
activity_logs	10000+	System activity audit trail
risk_alerts	50+	Active and resolved alerts

C. Autonomous Execution

The system operates autonomously through scheduled cron jobs:

```
// /api/cron/predictions
export async function GET(req: Request) {
  // Verify cron secret
  if (req.headers.get("authorization") !== `Bearer ${process.env.CRON_SECRET}`) {
    return new Response("Unauthorized", {
      status: 401
    });
  }

  // Fetch active customers
  const customers = await getActiveCustomers();

  // Run agent chain for high-risk customers
  for (const customer of customers) {
    const prediction = await runPrediction(customer);

    if (prediction.churn_probability > 0.7) {
      await agentOrchestrator
        .executeAgentChain(customer.id);
    }
  }

  return Response.json({ success: true });
}
```

2) *Anomaly Detection Cron (Every Hour):* Monitors KPIs and triggers Monitoring Agent when deviations are detected.

3) *Weekly Report Cron (Monday 9 AM):* Generates comprehensive executive summaries and delivers via email and Slack.

D. Agent Orchestration Implementation

The MultiAgentOrchestrator class coordinates agent execution:

```
async executeAgentChain(customerId: string) {
  // Step 1: Analyst Agent
  const analysis = await this.runAnalystAgent(
    customerId
  );

  // Step 2: Forecast Agent
  const forecast = await this.runForecastAgent(
    customerId, analysis
  );

  // Step 3: Decision Engine
  const decision = await this.runDecisionEngine(
    customerId, analysis, forecast
  );

  // Step 4: Action Executor
  if (decision.shouldAct) {
    await this.runActionAgent(decision, analysis);
  }
}
```

Each agent logs its decisions to the database with reasoning, confidence scores, and execution metadata.

E. Deployment

The system is deployed on Vercel with the following configuration:

- **Hosting:** Vercel Edge Network for global distribution
- **Cron Jobs:** Vercel Cron for scheduled tasks
- **Database:** Supabase hosted PostgreSQL
- **ML Service:** Deployed separately on Railway/Render

- **Environment:** Production environment variables for API keys

VI. EXPERIMENTAL RESULTS

A. Machine Learning Performance

TABLE II
CHURN PREDICTION MODEL PERFORMANCE

Metric	Value
Accuracy	86.94%
Precision	84.2%
Recall	88.7%
F1-Score	86.4%
AUC-ROC	0.91
Average Confidence	0.89




Fig. 6. Receiver Operating Characteristic (ROC) curve for the XGBoost churn prediction model.




Fig. 7. Relative feature importance generated by the XGBoost model.

1) Churn Prediction Model:

2) CLV Prediction Model:

- Mean Absolute Error (MAE): \$127
- Root Mean Squared Error (RMSE): \$189
- R^2 Score: 0.82

3) Anomaly Detection:

- True Positive Rate: 91.3%
- False Positive Rate: 4.2%
- Detection Latency: ≤ 1 second

B. Agent Performance

TABLE III
AGENT PERFORMANCE METRICS

Agent	Decisions	Confidence	Response Time
Monitoring	1247	89%	0.3s
Prediction	456	84%	2.1s
Decision	892	91%	0.5s
Action	324	98%	1.2s
Reporting	4521	99%	0.2s
Feedback	89	92%	5.0s

C. System Performance

1) Throughput:

- Predictions per hour: 600
- Agent decisions per day: 3,000+
- Actions executed per day: 150+
- Average end-to-end latency: 8.7 seconds

2) Reliability:

- System uptime: 99.7%
- Successful predictions: 99.2%
- Action execution success rate: 98.9%
- Database query success rate: 99.9%

3) Autonomy Metrics:

- Manual interventions per week: 2 (down from 20)
- Autonomous decision rate: 96%
- False positive alert rate: 3.8%
- Time to action (critical alerts): ≤ 2 minutes

D. Business Impact

1) Operational Efficiency:

- Reduction in manual monitoring: 90%
- Faster response to churn signals: 85% improvement
- Increased retention campaign effectiveness: 81% success rate
- Time saved per week: 15 hours

2) Cost Savings:

- Reduced customer churn: 12% improvement
- Prevented revenue loss: \$45K per month
- Operational cost reduction: 35%
- ROI: 340% in first 6 months

E. Comparison with Baseline

We compared SIM-OPS against a baseline manual monitoring system:

TABLE IV
COMPARISON WITH MANUAL BASELINE

Metric	Manual	SIM-OPS
Response Time	4-24 hours	12 minutes
Coverage	20%	100%
Accuracy	72%	87%
Cost per Alert	\$15	\$0.50
Human Hours/Week	20	2

VII. DISCUSSION

A. Key Findings

1) *Multi-Agent Coordination*: The six-agent architecture proved highly effective for decomposing complex business operations into specialized, manageable components. Agent communication through the message bus provided full traceability while maintaining loose coupling.

2) *Autonomous Operation*: The system successfully operates 24/7 with minimal human intervention. Scheduled cron jobs ensure continuous monitoring, and event-based triggers enable immediate response to critical situations.

3) *Explainability*: Every agent decision is logged with detailed reasoning, confidence scores, and execution metadata. This transparency is crucial for business stakeholders to trust autonomous systems.

B. Limitations

1) *Model Drift*: While the Feedback Agent monitors model performance, retraining requires manual intervention. Future work should implement fully automated retraining pipelines.

2) *Integration Complexity*: Setting up integrations with Slack, Jira, and email requires configuration and API keys. A more streamlined onboarding process would improve adoption.

3) *Scalability*: Current implementation processes customers sequentially. For larger customer bases (>10,000), parallel processing and distributed execution would be necessary.

4) *Cold Start Problem*: New deployments lack historical data for trend analysis. The system requires 7-14 days of operation to build sufficient history for accurate forecasting.

C. Lessons Learned

1) *Threshold Tuning*: Business rule thresholds require careful tuning based on industry and company size. We found that churn thresholds of 70% worked well for B2B SaaS but may need adjustment for other domains.

2) *Alert Fatigue*: Initial deployments generated too many medium-severity alerts. We implemented alert aggregation and daily summaries to reduce notification volume.

3) *Confidence Calibration*: ML model confidence scores needed calibration to align with business risk tolerance. We implemented confidence thresholds that balance false positives and false negatives.

VIII. FUTURE WORK

A. Enhanced ML Capabilities

- **Deep Learning Models**: Implement LSTM networks for time-series prediction
- **Natural Language Processing**: Analyze support ticket text for sentiment and urgency
- **Reinforcement Learning**: Optimize action selection based on historical outcomes
- **Automated Feature Engineering**: Use AutoML to discover new predictive features

B. Advanced Agent Capabilities

- **Negotiation Agent**: Coordinate with customers to prevent churn
- **Optimization Agent**: Recommend pricing and product changes
- **Simulation Agent**: Model "what-if" scenarios before taking actions
- **Learning Agent**: Continuously improve decision rules from outcomes

C. Scalability Improvements

- **Distributed Processing**: Use message queues (RabbitMQ, Kafka) for parallel execution
- **Caching Layer**: Implement Redis for frequently accessed predictions
- **Batch Optimization**: Process customers in optimized batches
- **Edge Computing**: Deploy prediction models closer to data sources

D. Integration Expansion

- **CRM Integration**: Salesforce, HubSpot for customer data
- **Analytics Platforms**: Mixpanel, Amplitude for behavioral data
- **Communication Tools**: Microsoft Teams, Discord for alerts
- **Data Warehouses**: Snowflake, BigQuery for historical analysis

E. Explainability and Trust

- **SHAP Values**: Implement model-agnostic explanations
- **Counterfactual Explanations**: Show what changes would alter predictions
- **Confidence Intervals**: Provide uncertainty quantification
- **Interactive Dashboards**: Allow stakeholders to explore decision reasoning

IX. CONCLUSION

This paper presented SIM-OPS, an autonomous multi-agent system for intelligent business operations management. Our six-agent architecture—comprising Monitoring, Prediction, Decision, Action, Reporting, and Feedback agents—demonstrates that complex business operations can be effectively automated through coordinated agent collaboration.

Key achievements include:

- 86.94% accuracy in churn prediction with 0.89 average confidence
- 90% reduction in manual intervention while maintaining 96% decision accuracy
- Sub-second response times for critical alerts with multi-channel delivery
- Full production deployment with 99.7% uptime and 24/7 autonomous operation
- Comprehensive audit trails and explainable decision-making

The system demonstrates that autonomous AI agents can successfully manage real-world business operations when designed with appropriate specialization, coordination mechanisms, and human oversight capabilities. Our open-source implementation provides a foundation for further research and practical deployment in various business domains.

Future work will focus on enhanced ML capabilities, advanced agent behaviors, scalability improvements, and broader integration with enterprise systems. We believe that autonomous multi-agent systems represent the future of business operations management, enabling organizations to respond faster, operate more efficiently, and make better data-driven decisions.

ACKNOWLEDGMENT

The authors would like to thank the open-source community for the excellent tools and libraries that made this work possible, including Next.js, FastAPI, Supabase, and the Python machine learning ecosystem.

REFERENCES

- [1] M. Wooldridge, *An Introduction to MultiAgent Systems*, 2nd ed. Chichester, UK: John Wiley & Sons, 2009.
- [2] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. Cambridge, MA: MIT Press, 2018.
- [3] P. Dayan and G. E. Hinton, "Feudal reinforcement learning," in *Advances in Neural Information Processing Systems*, 1992, pp. 271–278.
- [4] G. Shmueli and P. C. Koppius, "Predictive analytics in information systems research," *MIS Quarterly*, vol. 35, no. 3, pp. 553–572, 2016.
- [5] A. Bifet, R. Gavaldà, G. Holmes, and B. Pfahringer, *Machine Learning for Data Streams*. Cambridge, MA: MIT Press, 2018.
- [6] X. He, K. Zhao, and X. Chu, "AutoML: A survey of the state-of-the-art," *Knowledge-Based Systems*, vol. 212, p. 106622, 2021.
- [7] W. Verbeke, D. Martens, C. Mues, and B. Baesens, "Building comprehensible customer churn prediction models with advanced rule induction techniques," *Expert Systems with Applications*, vol. 38, no. 3, pp. 2354–2364, 2012.
- [8] Y. Xie, X. Li, E. W. T. Ngai, and W. Ying, "Customer churn prediction using improved balanced random forests," *Expert Systems with Applications*, vol. 36, no. 3, pp. 5445–5449, 2009.
- [9] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [10] Apache Airflow. (2015). *Apache Airflow Documentation*. [Online]. Available: <https://airflow.apache.org/>
- [11] Temporal Technologies. (2019). *Temporal Workflow Engine*. [Online]. Available: <https://temporal.io/>
- [12] M. Dumas, M. La Rosa, J. Mendling, and H. A. Reijers, *Fundamentals of Business Process Management*, 2nd ed. Berlin, Germany: Springer, 2018.
- [13] S. Yao et al., "ReAct: Synergizing Reasoning and Acting in Language Models," in *International Conference on Learning Representations (ICLR)*, 2023.